

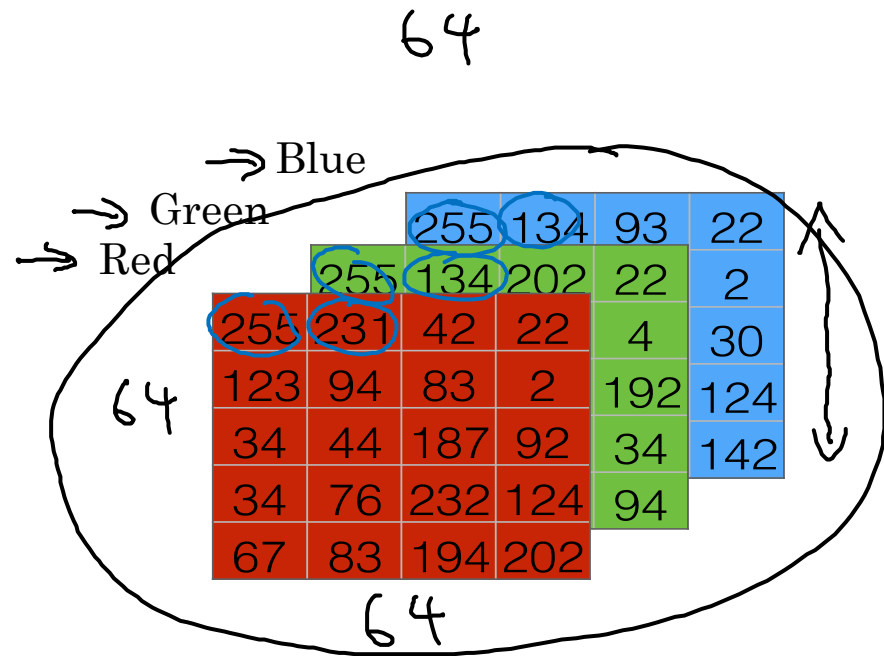



deeplearning.ai

Basics of Neural Network Programming

Binary Classification

Binary Classification



$X = \begin{bmatrix} 255 \\ 231 \\ \vdots \\ 255 \\ 134 \\ \vdots \end{bmatrix}$

$64 \times 64 \times 3 = 12288$

$n = n_x = 12288$

$X \rightarrow y$

Notation

$$(x, y) \quad x \in \mathbb{R}^{n_x}, y \in \{0, 1\}$$

$$m \text{ training examples: } \{(\underline{x}^{(1)}, \underline{y}^{(1)}), (\underline{x}^{(2)}, \underline{y}^{(2)}), \dots, (\underline{x}^{(m)}, \underline{y}^{(m)})\}$$

$$M = M_{\text{train}}$$


$$M_{\text{test}} = \# \text{ test examples.}$$

$$X = \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix}$$

$\xleftarrow{m} \quad \xrightarrow{\quad}$

$$X \in \mathbb{R}^{n_x \times m}$$

$$X.\text{shape} = (n_x, m)$$



dont do it like this
as it complicates
implementation

$$Y = [y^{(1)} \ y^{(2)} \ \dots \ y^{(m)}]$$

$$Y \in \mathbb{R}^{1 \times m}$$

$$Y.\text{shape} = (1, m)$$

.shape this is python command to find the shape of the matrix



deeplearning.ai

Basics of Neural Network Programming

Logistic Regression

Logistic Regression


Given x , want $\hat{y} = \frac{P(y=1|x)}{P(y=0|x) + P(y=1|x)}$

$$x \in \mathbb{R}^{n_x}$$

$$0 \leq \hat{y} \leq 1$$

Parameters: $\underline{w} \in \mathbb{R}^{n_x}$, $\underline{b} \in \mathbb{R}$.

Output $\hat{y} = \sigma(\underbrace{w^T x + b}_z)$



$$x_0 = 1, \quad x \in \mathbb{R}^{n_x + 1}$$
$$\hat{y} = \sigma(\theta^T x)$$

$$\Theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_{n_x} \end{bmatrix} \quad \left. \begin{matrix} \} b \leftarrow \\ \} w \leftarrow \end{matrix} \right\}$$

these are alternative denotation conventions that we are not going to follow here (he used this in other course i did)

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

If z large $\sigma(z) \approx \frac{1}{1 + 0} = 1$

If z large negative number

$$\sigma(z) = \frac{1}{1 + e^{-z}} \approx \frac{1}{1 + \text{Bignum}} \approx 0$$



deeplearning.ai

Basics of Neural Network Programming

Logistic Regression cost function

→ $\hat{y}^{(i)} = \sigma(w^T x^{(i)} + b)$, where $\sigma(z) = \frac{1}{1+e^{-z}}$ $z^{(i)} = w^T x^{(i)} + b$

Given $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$, want $\hat{y}^{(i)} \approx y^{(i)}$.


$x^{(i)}$
 $y^{(i)}$
 $z^{(i)}$ i -th example.

Loss (error) function:

$$\mathcal{L}(\hat{y}, y) = \frac{1}{2} (\hat{y} - y)^2$$

~~~~~

we don't use this since it's not convex and we have local minimum



$$\mathcal{L}(\hat{y}, y) = - (y \log \hat{y} + (1-y) \log (1-\hat{y})) \leftarrow$$

If  $y=1$ :  $\mathcal{L}(\hat{y}, y) = -\log \hat{y} \leftarrow$  Want  $\log \hat{y}$  large, want  $\hat{y}$  large · y here is sigmoid fct

If  $y=0$ :  $\mathcal{L}(\hat{y}, y) = -\log (1-\hat{y}) \leftarrow$  Want  $\log (1-\hat{y})$  large ... want  $\hat{y}$  small

Cost function:  $J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1-y^{(i)}) \log (1-\hat{y}^{(i)})]$



deeplearning.ai

# Basics of Neural Network Programming

---

## Logistic Regression cost function



# Logistic Regression cost function

$$\rightarrow \hat{y}^{(i)} = \sigma(w^T \underline{x}^{(i)} + b), \text{ where } \sigma(z^{(i)}) = \frac{1}{1+e^{-z^{(i)}}} \quad z^{(i)} = w^T x^{(i)} + b$$

Given  $\{(\underline{x}^{(1)}, \underline{y}^{(1)}), \dots, (\underline{x}^{(m)}, \underline{y}^{(m)})\}$ , want  $\hat{y}^{(i)} \approx y^{(i)}$ .

$x^{(i)}$   
 $y^{(i)}$   
 $z^{(i)}$   $i$ -th example.

Loss (error) function:

$$\mathcal{L}(\hat{y}, y) = \frac{1}{2} (\hat{y} - y)^2$$



$$\mathcal{L}(\hat{y}, y) = - (y \log \hat{y} + (1-y) \log (1-\hat{y})) \leftarrow$$

If  $y=1$ :  $\mathcal{L}(\hat{y}, y) = -\log \hat{y} \leftarrow$  Want  $\log \hat{y}$  large, want  $\hat{y}$  large.

If  $y=0$ :  $\mathcal{L}(\hat{y}, y) = -\log (1-\hat{y}) \leftarrow$  Want  $\log (1-\hat{y})$  large ... want  $\hat{y}$  small

$$\text{Cost function: } J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) = \frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1-y^{(i)}) \log (1-\hat{y}^{(i)})]$$

Andrew Ng



deeplearning.ai

# Basics of Neural Network Programming

---



## Gradient Descent

# Gradient Descent


Recap:  $\hat{y} = \sigma(w^T x + b)$ ,  $\sigma(z) = \frac{1}{1+e^{-z}}$   $\leftarrow$

$$\underline{J(w, b)} = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\underline{\hat{y}^{(i)}} , \underline{y^{(i)}}) = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$

Want to find  $w, b$  that minimize  $J(w, b)$



# Gradient Descent



Repeat {  
 $w := w - \alpha \frac{dJ(w)}{dw}$   
}

learning rate (pointing to  $\alpha$ )

"dw" (pointing to  $\frac{dJ(w)}{dw}$ )

$w := w - \alpha dw$

$\frac{dJ(w)}{dw} = ?$  We need to find these partial derivatives.

$J(w, b)$

$$w := w - \alpha \frac{\partial J(w, b)}{\partial w}$$

$$b := b - \alpha \frac{\partial J(w, b)}{\partial b}$$

$\frac{\partial J(w, b)}{\partial w}$  (in a red box)

$\frac{\partial J(w, b)}{\partial b}$  (in a red box)

$\partial$  (circled in blue, with an arrow pointing to it from the text "partial derivative")

$J$  (circled in blue, with an arrow pointing to it from the text "partial derivative")

$dw$  (in red, with an arrow pointing to it from the red box containing  $\frac{\partial J(w, b)}{\partial w}$ )

$db$  (in red, with an arrow pointing to it from the red box containing  $\frac{\partial J(w, b)}{\partial b}$ )




deeplearning.ai

# Basics of Neural Network Programming

---

## Derivatives

# Intuition about derivatives



$$\frac{0.003}{0.001} \quad \text{height} \\ \text{width}$$

$\rightarrow a = 2 \quad f(a) = 6$   
 $a = 2.001 \quad f(a) = 6.003$   
 slope (derivative) of  $f(a)$   
 at  $a = 2$  is 3

$\rightarrow a = 5 \quad f(a) = 15$   
 $a = 5.001 \quad f(a) = 15.003$   
 slope at  $a = 5$  is also 3

$$\frac{df(a)}{da} = 3 = \frac{d}{da} f(a)$$

$0.001 \leftarrow$   
 $0.000000001$   
 $0.0000000001$




deeplearning.ai

# Basics of Neural Network Programming

---

More derivatives  
examples

# Intuition about derivatives



$$f(a) = a^2$$

$$\frac{\text{height}}{\text{width}}$$

$$\frac{d}{da} a^2 = 2a$$

$$0.001$$

$$(2a) \times 0.001$$

$a = 2$                        $f(a) = 4$   
 $a = 2.001$                  $f(a) \approx 4.004$   
                                           $(4.004 \text{ } \boxed{004})$

slope (derivative) of  $f(a)$  at  $a=2$  is 4.

$$\frac{d}{da} f(a) = 4 \quad \text{when } a=2$$

$$\begin{array}{ll} a=5 & f(a)=25 \\ a=5.001 & f(a) \approx 25.\underline{010} \end{array}$$

$$\frac{d}{da} f(a) = 10 \quad \text{when} \quad a = 5$$

$$\frac{d}{da} f(a) = \frac{d}{da} a^2 = \boxed{2a}$$



# More derivative examples

$$f(a) = a^2$$

$$\frac{d}{da} f(a) = \frac{2a}{4}$$

$$a = 2$$

$$f(a) = 4$$

$$a = 2.001$$

$$f(a) \approx 4.004$$

$$f(a) = a^3$$

$$\frac{d}{da} f(a) = \frac{3a^2}{3 \times 2^2 = 12}$$

$$a = 2$$

$$f(a) = 8$$


$$a = \underline{2.001}$$

$$f(a) \approx \underline{8.012}$$

$$f(a) = \log_e(a)$$
  

$$\ln(a)$$

$$\frac{d}{da} f(a) = \frac{1}{a}$$



$$\frac{d}{da} f(a) = \frac{1}{2}$$

$$a = 2$$

$$f(a) \approx 0.69315$$

$$\downarrow$$
  

$$a = \underline{2.001}$$

$$\downarrow$$
  

$$\underline{f(a) \approx 0.69365}$$

$$\downarrow$$
  

$$0.0005$$
  

$$\swarrow$$
  

$$\underline{0.0005}$$



deeplearning.ai

# Basics of Neural Network Programming


---

## Computation Graph

# Computation Graph

$$J(a,b,c) = 3(a + \underbrace{bc}_u) = 3(5 + \underbrace{3 \times 2}_v) = 33$$

$$\begin{aligned} u &= bc \\ v &= a + u \\ J &= 3v \end{aligned}$$



So, the computation graph comes in handy when there is some distinguished or some special output variable, such as  $J$  in this case, that you want to optimize. And in the case of a logistic regression,  $J$  is of course the cost function that we're trying to minimize. And what we're seeing in this little example is that, through a left-to-right pass, you can compute the value of  $J$ . And what we'll see in the next couple of slides is that in order to compute derivatives, there'll be a right-to-left pass like this, kind of going in the opposite direction as the blue arrows. That would be most natural for computing the derivatives. So to recap, the computation graph organizes a computation with this blue arrow, left-to-right computation. Let's refer to the next video how you can do the backward red arrow right-to-left computation of the derivatives. Let's go on to the next video.

Andrew Ng

Andrew Ng



deeplearning.ai


# Basics of Neural Network Programming

---

## Derivatives with a Computation Graph



# Computing derivatives



$$\frac{dJ}{dv} = ? = 3$$

$a \rightarrow v \rightarrow J$

$$J = 3v$$

$$v = 11 \rightarrow 11.001$$

$$J = 33 \rightarrow 33.003$$

$$f(a) = 3a$$

$$\frac{df(a)}{da} = \frac{df}{da} = 3$$

$$J = 3v$$

$$\frac{dJ}{dv} = 3$$

$$\frac{dJ}{da} = 3 = \frac{dJ}{dv} \frac{dv}{da}$$

$$\frac{dv}{da} = 1$$

$$\frac{\partial \text{Final Output Var}}{\partial \text{var}}$$

$$a = 5 \rightarrow 5.001$$

$$\rightarrow v = 11 \rightarrow 11.001$$

$$J = 33 \rightarrow 33.003$$


$$\frac{dJ}{d\text{var}}$$

"dvar"

Andrew Ng

Andrew Ng

# Computing derivatives



$$\underline{\frac{dJ}{du}} = 3 = \underbrace{\frac{dJ}{dv}}_3 \cdot \underbrace{\frac{dv}{du}}_1$$

$$\frac{dJ}{db} = \frac{dJ}{du} \cdot \frac{du}{db} = 6$$

$\underbrace{\frac{dJ}{du}}_{\rightarrow 3} \quad \underbrace{\frac{du}{db}}_{=2}$

$$\frac{dJ}{dc} = \frac{dJ}{du} \cdot \frac{du}{dc} = 9$$

$\underbrace{\frac{dJ}{du}}_{\rightarrow 3} \quad \underbrace{\frac{du}{dc}}_{=3}$

$$\begin{aligned} u = 6 &\rightarrow 6.001 \\ v = 11 &\rightarrow 11.001 \\ J = 33 &\rightarrow 33.003 \end{aligned}$$

$$b = 3 \rightarrow 3.001$$

$$\begin{aligned} u = b \cdot c = 6 &\rightarrow 6.002 \\ J = 33.006 \end{aligned}$$

$$\begin{aligned} v = 11.002 \\ J = 3v \end{aligned}$$

$$\begin{aligned} c = 2 \\ 1006 \end{aligned}$$

Andrew Ng





deeplearning.ai

# Basics of Neural Network Programming

---

## Logistic Regression Gradient descent

thsi si for one training example

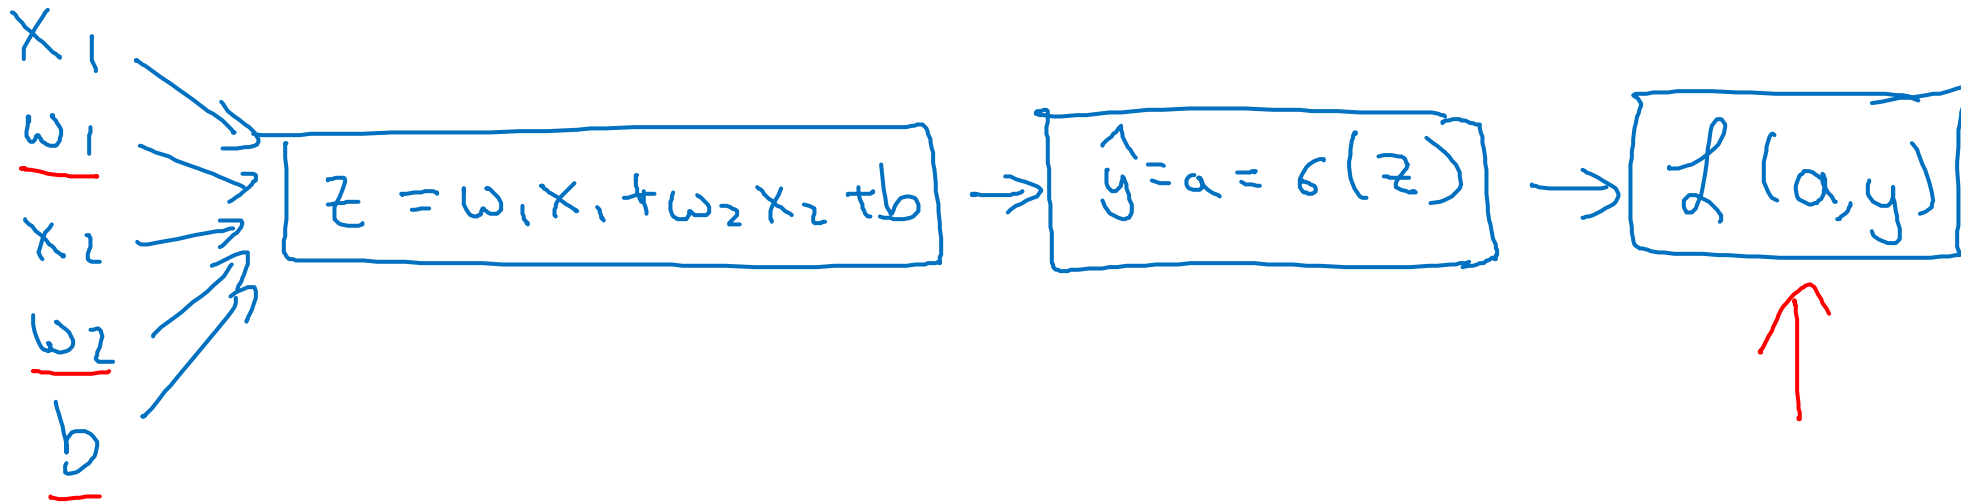


# Logistic regression recap


→  $z = w^T x + b$

→  $\hat{y} = a = \sigma(\underline{z})$

→  $\mathcal{L}(a, y) = -(y \log(a) + (1 - y) \log(1 - a))$



# Logistic regression derivatives





deeplearning.ai

# Basics of Neural Network Programming

---

Gradient descent  
on *m* examples

# Logistic regression on $m$ examples

we just using this training example

$$\underline{J(w, b)} = \frac{1}{m} \sum_{i=1}^m \ell(a^{(i)}, y^{(i)})$$
$$\rightarrow a^{(i)} = \hat{y}^{(i)} = \sigma(z^{(i)}) = \sigma(w^T x^{(i)} + b)$$

$$(x^{(i)}, y^{(i)})$$

$$\underline{dw_1^{(i)}}, \underline{dw_2^{(i)}}, \underline{db^{(i)}}$$

$$\underline{\frac{\partial}{\partial w_1} J(w, b)} = \frac{1}{m} \sum_{i=1}^m \underbrace{\frac{\partial}{\partial w_1} \ell(a^{(i)}, y^{(i)})}_{\underline{dw_1^{(i)}} \rightarrow (x^{(i)}, y^{(i)})}$$

we showed in previous slide how to compute this for a single training sample

So for  $m$  training examples we just compute those partial derivatives for all the samples and do an average as is done above.

# Logistic regression on $m$ examples

$$J=0; \underline{dw_1}=0; \underline{dw_2}=0; \underline{db}=0$$

initialize them a those values.

→ For  $i=1$  to  $m$  this is for the training sets, from  $i$  to  $m$

$$z^{(i)} = w^T x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$J += -[y^{(i)} \log a^{(i)} + (1-y^{(i)}) \log(1-a^{(i)})]$$

$$\underline{dz^{(i)}} = a^{(i)} - y^{(i)}$$

$$dw_1 += x_1^{(i)} dz^{(i)}$$

$$dw_2 += x_2^{(i)} dz^{(i)}$$

$$db += dz^{(i)}$$

$$\underline{n=2}$$

we assume we have just 2 features, in reality they will be a lot.

$dw_3$   
 $\vdots$   
 $dw_n$

$$J /= m \leftarrow$$

You still need to divide by  $n$  because we are computing averages.

$$dw_1 /= m; dw_2 /= m; db /= m. \leftarrow$$

We do average, check previews slide why.

$$dw_1 = \frac{\partial J}{\partial w_1}$$

$$w_1 := w_1 - \alpha \underline{dw_1}$$

$$w_2 := w_2 - \alpha \underline{dw_2}$$

$$b := b - \alpha \underline{db}$$

Vectorization



deeplearning.ai

# Basics of Neural Network Programming

---

## Vectorization

# What is vectorization?

$$z = \underbrace{w^T x}_{\text{dot product}} + b$$

$$w = \begin{bmatrix} \vdots \\ \vdots \\ \vdots \end{bmatrix} \quad x = \begin{bmatrix} \vdots \\ \vdots \\ \vdots \end{bmatrix}$$

$$w \in \mathbb{R}^{n_x}$$

$$x \in \mathbb{R}^{n_x}$$

Non-vectorized:

$$z = 0$$

for  $i$  in  $\text{range}(n-x)$ :

$$z += w[i] * x[i]$$

$$z += b$$

Vectorized

$$z = \underbrace{\text{np.dot}(w, x)}_{w^T x} + b$$

$\Rightarrow$  GPU } SIMD - single instruction  
 $\Rightarrow$  CPU } multiple data.



deeplearning.ai

# Basics of Neural Network Programming

---

## More vectorization examples



# Neural network programming guideline

Whenever possible, avoid explicit for-loops.

# Neural network programming guideline

Whenever possible, avoid explicit for-loops.

$$u = Av$$

$$u_i = \sum_j A_{ij} v_j$$

$$u = \text{np.zeros}(n, 1)$$

for i ... ←

for j ... ←

$$u[i] += A[i][j] * v[j]$$

$$u = \text{np.dot}(A, v)$$

# Vectors and matrix valued functions

Say you need to apply the exponential operation on every element of a matrix/vector.

$$v = \begin{bmatrix} v_1 \\ \vdots \\ v_n \end{bmatrix} \rightarrow u = \begin{bmatrix} e^{v_1} \\ e^{v_2} \\ \vdots \\ e^{v_n} \end{bmatrix}$$

```
→ u = np.zeros((n,1))  
→ for i in range(n):  
    → u[i]=math.exp(v[i])
```

```
import numpy as np  
u = np.exp(v)  
  
np.log(v)  
np.abs(v)  
np.maximum(v, 0)  
v**2  
1/v
```

# Logistic regression derivatives

$$J = 0, \quad \boxed{\cancel{dw1 = 0, dw2 = 0}}, \quad db = 0$$

$$dw = np.zeros((n-x, 1))$$

initialize a vector of zeros with nx elements( the number of weights)

→ for i = 1 to n:

$$z^{(i)} = w^T x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$J += -[y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

$$dz^{(i)} = a^{(i)}(1 - a^{(i)})$$

$$\cancel{dw_1 += x_1^{(i)} dz^{(i)}}$$

$$\cancel{dw_2 += x_2^{(i)} dz^{(i)}}$$

$$db += dz^{(i)}$$

$n_x = 2$

$$dw += x^{(i)} dz^{(i)}$$

$$J = J/m, \quad \boxed{\cancel{dw_1 = dw_1/m, dw_2 = dw_2/m}}, \quad db = db/m$$

$$dw /= m.$$



deeplearning.ai

# Basics of Neural Network Programming

---

## Vectorizing Logistic Regression

Here we vectorise implementation without using any loop

# Vectorizing Logistic Regression

$$\begin{aligned} \rightarrow z^{(1)} &= w^T x^{(1)} + b \\ \rightarrow a^{(1)} &= \sigma(z^{(1)}) \end{aligned} \quad \begin{aligned} z^{(2)} &= w^T x^{(2)} + b \\ a^{(2)} &= \sigma(z^{(2)}) \end{aligned} \quad \begin{aligned} z^{(3)} &= w^T x^{(3)} + b \\ a^{(3)} &= \sigma(z^{(3)}) \end{aligned}$$

$$\underline{X} = \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix}$$

↑

$$\begin{matrix} (n_x, m) \\ \mathbb{R}^{n_x \times m} \end{matrix}$$

$$\underline{w}^T \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix}$$

$$\underline{z} = \begin{bmatrix} z^{(1)} & z^{(2)} & \dots & z^{(m)} \end{bmatrix} = \underline{w}^T \underline{X} + \begin{bmatrix} b & b & \dots & b \end{bmatrix}_{1 \times m} = \begin{bmatrix} w^T x^{(1)} + b & w^T x^{(2)} + b & \dots & w^T x^{(m)} + b \end{bmatrix}_{1 \times m}$$

$$\rightarrow \underline{z} = \text{np.dot}(\underline{w.T}, \underline{X}) + \underline{b}$$

$$\underline{A} = \begin{bmatrix} a^{(1)} & a^{(2)} & \dots & a^{(m)} \end{bmatrix} = \sigma(\underline{z})$$

"Broadcasting"

b here is just real number but python expands it to a row vector b like above. Its called BROADCASTING

SEE DETAIL OF THIS IN PROG ASSIGNMENT.

Andrew Ng



deeplearning.ai

# Basics of Neural Network Programming

---

## Vectorizing Logistic Regression's Gradient Computation

# Vectorizing Logistic Regression

$$dz^{(1)} = a^{(1)} - y^{(1)}$$

$$dz^{(2)} = a^{(2)} - y^{(2)}$$

.....

$$dz = [dz^{(1)} \quad dz^{(2)} \quad \dots \quad dz^{(m)}]$$

$1 \times m$

←

$$A = [a^{(1)} \quad \dots \quad a^{(m)}]$$

$$Y = [y^{(1)} \quad \dots \quad y^{(m)}]$$

$$\rightarrow dz = A - Y = [a^{(1)} - y^{(1)} \quad a^{(2)} - y^{(2)} \quad \dots]$$

this is how we compute dz in vectorised form

$$\rightarrow dw = 0$$

$$dw += \frac{x^{(1)} dz^{(1)}}{m}$$

$$dw += \frac{x^{(2)} dz^{(2)}}{m}$$

⋮

$$dw /= m$$

$$db = 0$$

$$db += dz^{(1)}$$

$$db += dz^{(2)}$$

$$\vdots$$

$$db += dz^{(m)}$$

$$db /= m$$

$$db = \frac{1}{m} \sum_{i=1}^m dz^{(i)}$$

$$= \frac{1}{m} \text{np.sum}(dz)$$

$$dw = \frac{1}{m} X dz^T$$

$$= \frac{1}{m} \begin{bmatrix} x^{(1)} & \dots & x^{(m)} \\ 1 & & 1 \end{bmatrix} \begin{bmatrix} dz^{(1)} \\ \vdots \\ dz^{(m)} \end{bmatrix}$$

$$= \frac{1}{m} \left[ \underbrace{x^{(1)} dz^{(1)} + \dots + x^{(m)} dz^{(m)}}_{n \times 1} \right]$$

this is a vector of all the X input values



# Implementing Logistic Regression

NO WAY TO GET  
RID OF THIS  
LOOP FOR NR OF  
ITERATIONS

THIS IS OUR NON EFFICIENT IMPLEMENTATION

$J = 0, dw_1 = 0, dw_2 = 0, db = 0$

for  $i = 1$  to  $m$ :

WE NOW GOT RID OF THIS ONE ALSO

$$z^{(i)} = w^T x^{(i)} + b \leftarrow$$

$$a^{(i)} = \sigma(z^{(i)}) \leftarrow$$

$$J += -[y^{(i)} \log a^{(i)} + (1 - y^{(i)}) \log(1 - a^{(i)})]$$

$$dz^{(i)} = a^{(i)} - y^{(i)} \leftarrow$$

$$\left[ \begin{array}{l} dw_1 += x_1^{(i)} dz^{(i)} \\ dw_2 += x_2^{(i)} dz^{(i)} \\ db += dz^{(i)} \end{array} \right] \quad dw += x^{(i)} * dz^{(i)}$$

we got rid of a potential  
loop here.

$$J = J/m, dw_1 = dw_1/m, dw_2 = dw_2/m$$

$$db = db/m$$

for iter in range(1000):

$$z = w^T X + b$$
$$= np.dot(w.T, X) + b$$

$$A = \sigma(z)$$

$$dz = A - Y$$

$$dw = \frac{1}{m} X dz^T$$

$$db = \frac{1}{m} np.sum(dz)$$

$$w := w - \alpha dw$$

$$b := b - \alpha db$$

FINALD CODE HOW TO  
IMPLEMENT LOGISTIC REGRESSION



deeplearning.ai

# Basics of Neural Network Programming

---

## Broadcasting in Python

# Broadcasting example

Calories from Carbs, Proteins, Fats in 100g of different foods:

|         | Apples | Beef  | Eggs | Potatoes |
|---------|--------|-------|------|----------|
| Carb    | 56.0   | 0.0   | 4.4  | 68.0     |
| Protein | 1.2    | 104.0 | 52.0 | 8.0      |
| Fat     | 1.8    | 135.0 | 99.0 | 0.9      |

= A  
(3,4)

59 cal  
 $\frac{56}{59} \approx 94.9\%$

Calculate % of calories from Carb, Protein, Fat. Can you do this without explicit for-loop?

```
cal = A.sum(axis = 0)
percentage = 100*A/(cal.reshape(1,4))
```

↑(3,4) / (1,4)

```
>>> import numpy as np
>>> A = np.array([[56.0,0.0,4.4,68.0],
[1.2,104.0,52.0,8.0],[1.8,135.0,99.0,0.9]])
>>> A
array([[ 56. ,  0. ,  4.4, 68. ],
[  1.2, 104. , 52. ,  8. ],
[  1.8, 135. , 99. ,  0.9]])
>>> cal = A.sum(axis = 0)
>>> cal
array([ 59. , 239. , 155.4, 76.9])
>>> cal.reshape(1,4)
array([[ 59. , 239. , 155.4, 76.9]])
>>> percentage = 100*A/cal
>>> percentage
array([[94.91525424,  0. ,  2.83140283,
88.42652796],
[ 2.03389831, 43.51464435, 33.46203346,
10.40312094],
[ 3.05084746, 56.48535565, 63.70656371,
1.17035111]])
>>>
```

# Broadcasting example

$$\begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix} + \begin{bmatrix} 100 \\ 100 \\ 100 \\ 100 \end{bmatrix} \quad \text{100}$$

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} + \begin{bmatrix} 100 & 200 & 300 \\ 100 & 200 & 300 \end{bmatrix}$$

$(m,n) \quad (2,3) \quad (1,n) \rightsquigarrow (m,n) \quad (2,3)$

↓   ↓   ↓

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} + \begin{bmatrix} 100 & 100 & 100 \\ 200 & 200 & 200 \end{bmatrix} =$$

$(m,n) \quad (m,1) \rightsquigarrow (m,n)$

←   ←

# General Principle

$$\begin{array}{ccc} (m, n) & + & (1, n) \\ \text{matrix} & \times & \rightsquigarrow (m, n) \\ \hline & / & \end{array}$$
$$(m, 1) \rightsquigarrow (m, n)$$

$$\begin{array}{ccc} (m, 1) & + & \mathbb{R} \\ \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} & + & 100 \\ [1 \ 2 \ 3] & + & 100 \end{array} = \begin{bmatrix} 101 \\ 102 \\ 103 \end{bmatrix}$$
$$= [101 \quad 102 \quad 103]$$

Matlab/Octave: bsxfun



deeplearning.ai

# Basics of Neural Network Programming

---

A note on python/  
numpy vectors

# Python Demo

# Python / numpy vectors

```
import numpy as np

a = np.random.randn(5)

a = np.random.randn(5, 1)

a = np.random.randn(1, 5)

assert(a.shape == (5, 1))
```





deeplearning.ai

# Basics of Neural Network Programming

---

Quick tour of Jupyter/  
ipython notebooks

Logistic Regression with a Neural Network mindset - Student 06-16

Last Checkpoint: 06/24/2017 (unsaved changes)

Control Panel

Logout

File

Edit

View

Insert

Cell

Kernel

Widgets

Help

Python 3

+

3

↺

↻

↱

↲

↵

⏏

⏏

Code

Cell Toolbar

Submit Assignment

## 1) Packages

Let's first import all the packages that you will need during this assignment.

- [numpy](#) is the fundamental package for scientific computing with Python.
- [h5py](#) is a common package to interact with a dataset that is stored on an H5 file.
- [matplotlib](#) is a famous library to plot graphs in Python.

```
In [6]: import numpy as np
import matplotlib.pyplot as plt
import h5py

%matplotlib inline

### START CODE HERE ### (= 3 lines of code)
print("Hello World")
### END CODE HERE ###
```

Hello World

## 2) Overview of the Problem set

**Problem Statement:** You are given a dataset ("data.h5") containing:

- a training set of  $n_{\text{train}}$  images labeled as cat (1) or non-cat (0)
- a test set of  $n_{\text{test}}$  images labeled as cat and non-cat



deeplearning.ai

# Basics of Neural Network Programming

---

Explanation of logistic  
regression cost function  
(Optional)

# Logistic regression cost function

$$\begin{aligned} \rightarrow & \text{If } y = 1: p(y|x) = \hat{y} \\ \rightarrow & \text{If } y = 0: p(y|x) = 1 - \hat{y} \end{aligned} \quad \left. \vphantom{\begin{aligned} \rightarrow & \text{If } y = 1: p(y|x) = \hat{y} \\ \rightarrow & \text{If } y = 0: p(y|x) = 1 - \hat{y} \end{aligned}} \right\} p(y|x)$$

$$p(y|x) = \hat{y}^y (1-\hat{y})^{(1-y)} \quad \leftarrow$$

$$\text{If } y=1: p(y|x) = \hat{y} \underbrace{(1-\hat{y})^0}_{=1}$$

$$\text{If } y=0: p(y|x) = \underbrace{\hat{y}^0}_1 (1-\hat{y})^{(1-0)} = 1 \times (1-\hat{y}) = \underline{1-\hat{y}}$$

$$\begin{aligned} \log p(y|x) &= \log \hat{y}^y (1-\hat{y})^{(1-y)} = y \log \hat{y} + (1-y) \log (1-\hat{y}) \\ &= -\frac{1}{\lambda} \underbrace{f(\hat{y}, y)} \end{aligned}$$

Andrew Ng

# Cost on $m$ examples

$$\log p(\text{labels in training set}) = \log \prod_{i=1}^m p(y^{(i)} | x^{(i)}) \leftarrow$$

$$\log p(\text{---}) = \sum_{i=1}^m \underbrace{\log p(y^{(i)} | x^{(i)})}_{= \mathcal{L}(\hat{y}^{(i)}, y^{(i)})}$$

Maximum likelihood  
estimator  $\nwarrow$

$$= - \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

$$\text{Cost:} \quad \underbrace{J(w, b)}_{\text{(minimize)}} = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$