



deeplearning.ai

# Setting up your ML application

---

## Train/dev/test sets

# Applied ML is a highly iterative process


# layers

## # hidden units

# learning rates

# activation functions

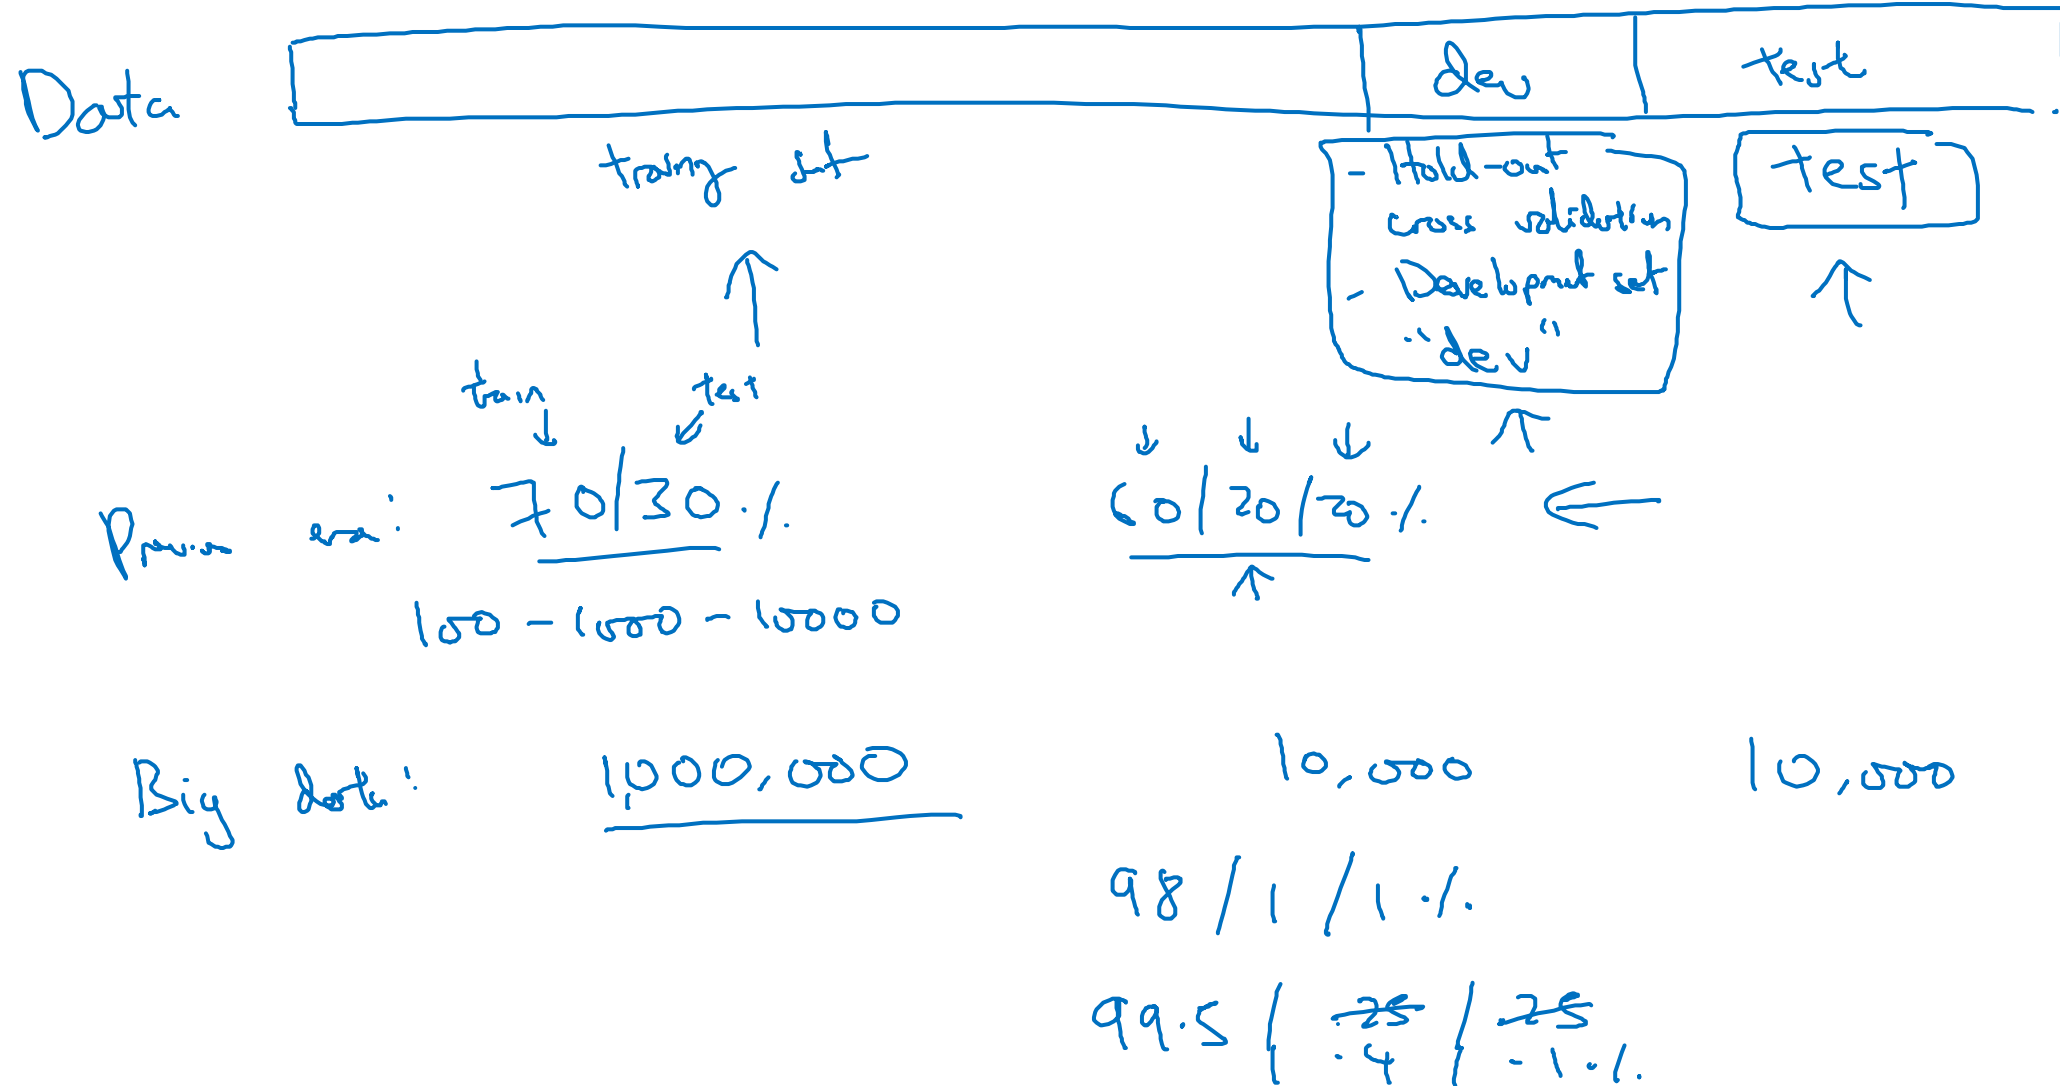
...



NLP, Vision, Speech, Structured data


```
graph LR; A["NLP, Vision, Speech"] --> B[Ads]; C[Structured data] --- D[Search]; C --- E[Security]; C --- F["logistic ..."]; E --> G[ ]; F --> G;
```

# Train/dev/test sets





# Mismatched train/test distribution

Certs



→ Make sure dev and test come from same distribution.



Not having a test set might be okay. (Only dev set.)






deeplearning.ai

# Setting up your ML application

---


## Bias/Variance

# Bias and Variance



# Bias and Variance


## Cat classification



Train set error:	1%	15% ←	15%	0.5%
Dev set error:	11%	16% ←	30%	1%
	high variance ↑	high bias ↑ ↑	high bias & high variance	low bias low variance ↑
Human: ~0%				
Optimal (Bayes) error: ~0% 15%				

Blurry images

# High bias and high variance



high bias  
high variance





deeplearning.ai

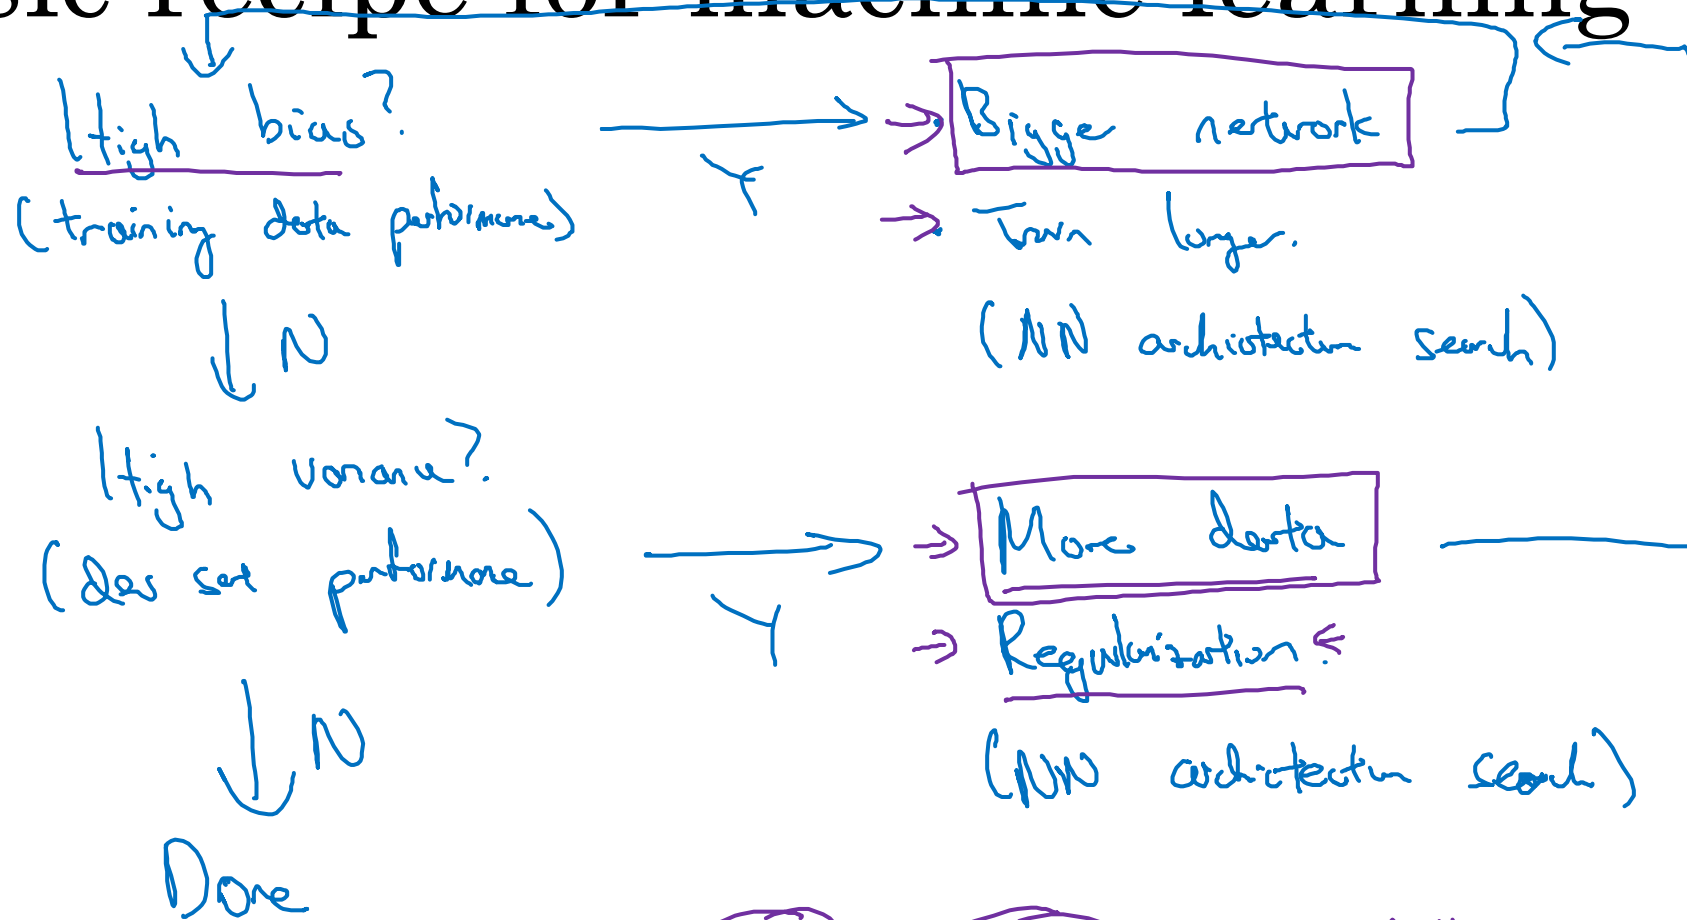
# Setting up your ML application

---

## Basic “recipe” for machine learning

# Basic “recipe” for machine learning

# Basic recipe for machine learning





**deeplearning.ai**

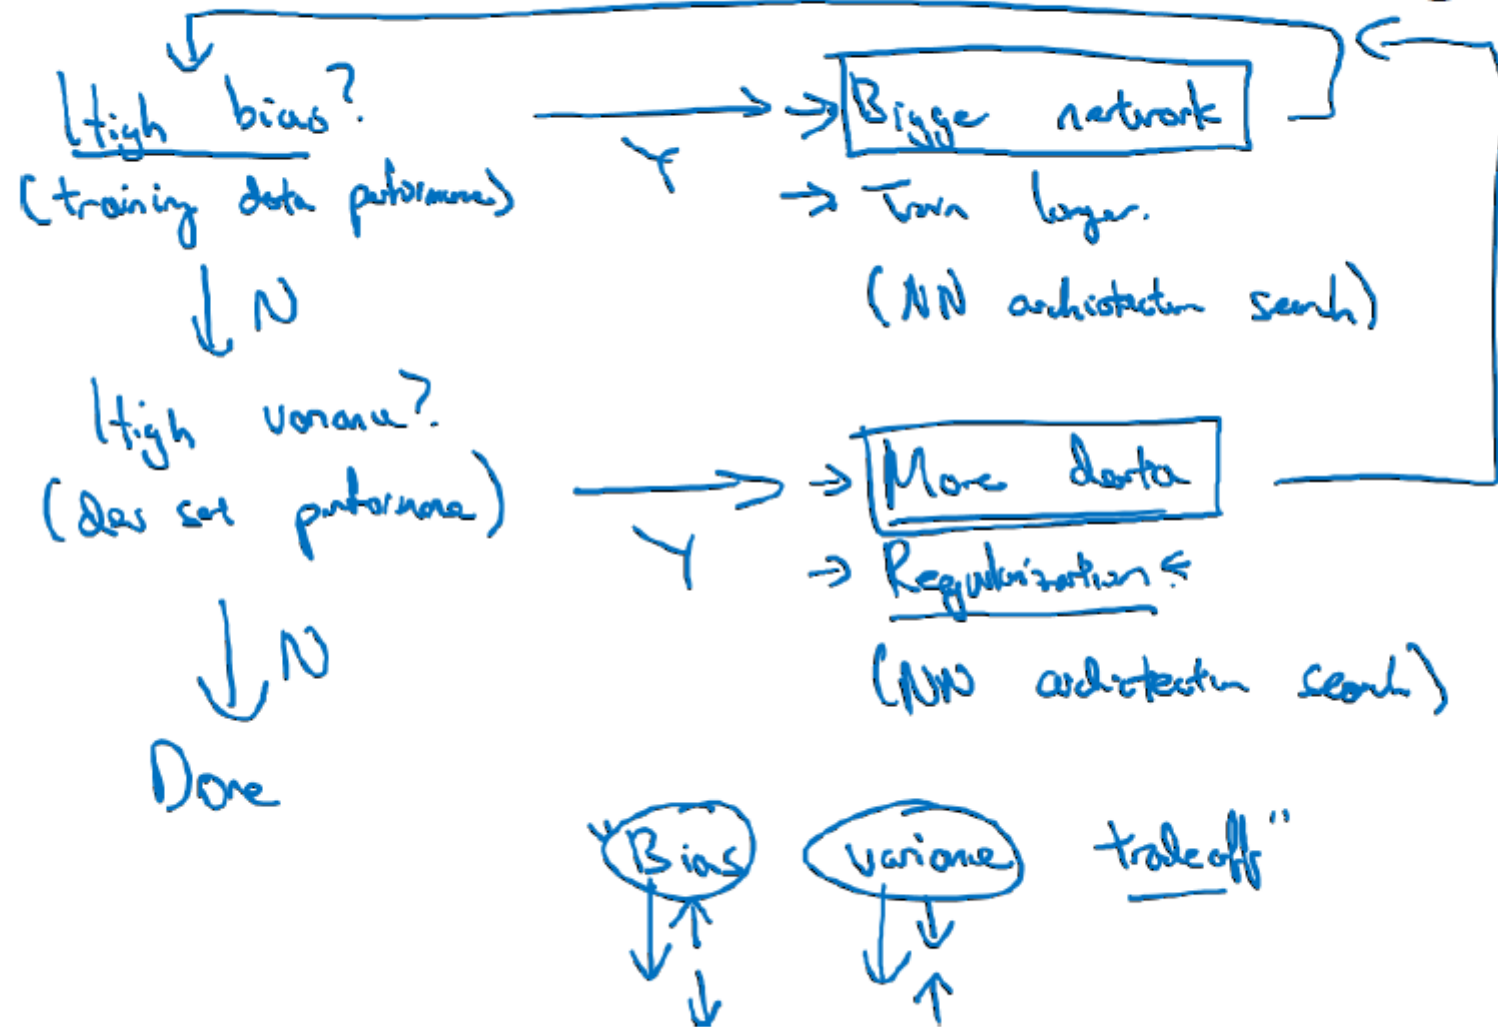
# Setting up your ML application

---

## Basic “recipe” for machine learning

# Basic “recipe” for machine learning

# Basic recipe for machine learning





deeplearning.ai

# Regularizing your neural network

---

## Regularization

# Logistic regression

$$\underline{w \in \mathbb{R}^{n_x}}, \underline{b \in \mathbb{R}}$$

$\lambda$  = regularization parameter  
lambda lambda

$$\min_{w,b} J(w,b)$$

$$J(w,b) = \underbrace{\frac{1}{m} \sum_{i=1}^m \ell(y^{(i)}, \hat{y}^{(i)})}_{\text{loss}} + \frac{\lambda}{2m} \underbrace{\|w\|_2^2}_{\text{L2 regularization}}$$

~~$+\frac{\lambda}{2m} b^2$~~   
omit

$L_2$  regularization  $\underline{\|w\|_2^2} = \sum_{j=1}^{n_x} w_j^2 = w^T w \leftarrow$

$L_1$  regularization  $\frac{\lambda}{2m} \sum_{j=1}^{n_x} |w_j| = \frac{\lambda}{2m} \|w\|_1$

$w$  will be sparse



# Neural network

$$\rightarrow J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = \underbrace{\frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, \hat{y}^{(i)})}_{\text{loss}} + \underbrace{\frac{\lambda}{2n} \sum_{l=1}^L \|w^{[l]}\|_F^2}_{\text{weight decay}}$$

$$\|w^{[l]}\|_F^2 = \sum_{i=1}^{n^{[l]}} \sum_{j=1}^{n^{[l-1]}} (w_{ij}^{[l]})^2$$

$w^{[l]}: \begin{matrix} n^{[l]} & n^{[l-1]} \\ \uparrow & \uparrow \end{matrix}$

"Frobenius norm"

$\|\cdot\|_2^2$

$\|\cdot\|_F^2$

$$dw^{[l]} = \left[ (\text{from backprop}) + \frac{\lambda}{n} w^{[l]} \right]$$

$$\rightarrow w^{[l]} := w^{[l]} - \alpha dw^{[l]}$$

$$\frac{\partial J}{\partial w^{[l]}} = dw^{[l]}$$

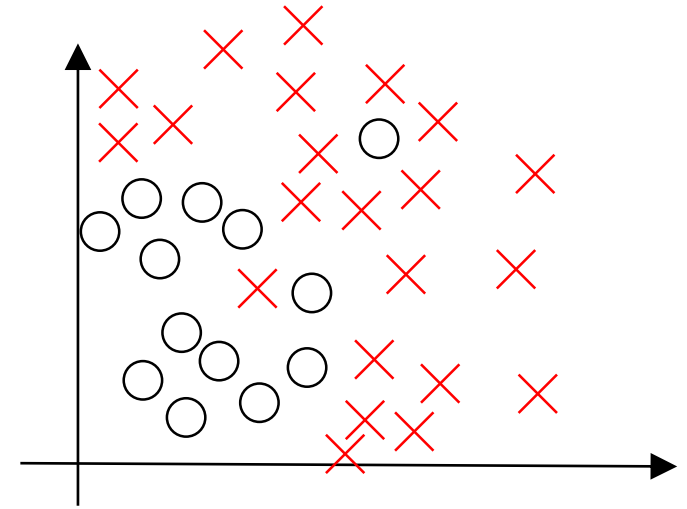
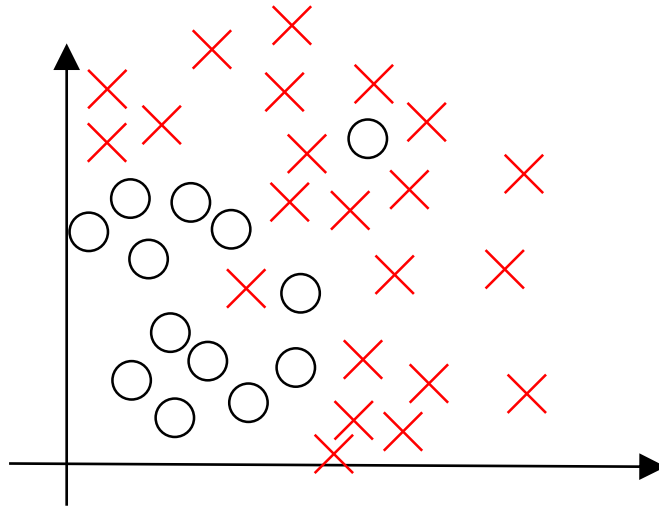
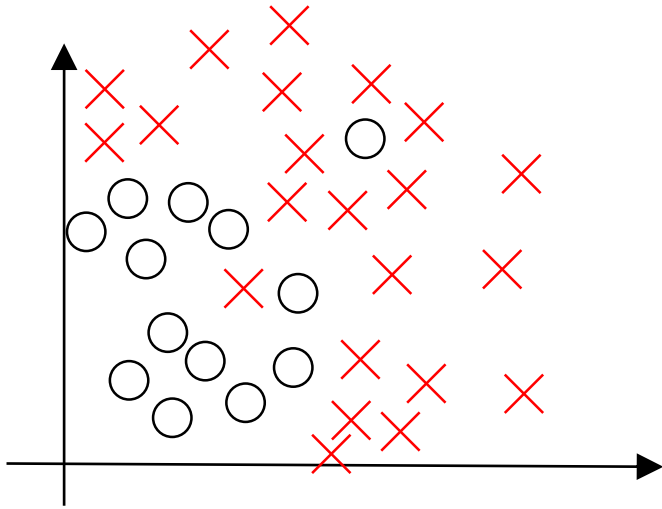

"Weight decay"

$$w^{[l]} := w^{[l]} - \alpha \left[ (\text{from backprop}) + \frac{\lambda}{n} w^{[l]} \right]$$

$$= w^{[l]} - \frac{\alpha \lambda}{n} w^{[l]} - \alpha (\text{from backprop})$$

$$= \underbrace{\left(1 - \frac{\alpha \lambda}{n}\right)}_{\leq 1} \underbrace{w^{[l]}}_{\text{weight}} - \alpha (\text{from backprop})$$

# How does regularization prevent overfitting?



# How does regularization prevent overfitting?




deeplearning.ai

# Regularizing your neural network

---




## Why regularization reduces overfitting

# How does regularization prevent overfitting?




$$J(\mathbf{w}^{(1)}, \mathbf{b}^{(1)}) = \frac{1}{n} \sum_{i=1}^n \ell(y^{(i)}, \hat{y}^{(i)}) + \frac{\lambda}{2n} \sum_{l=1}^L \underbrace{\|\mathbf{w}^{(l)}\|_F^2}_{\text{Frobenius norm squared}}$$

$$\mathbf{w}^{(1)} \approx 0$$



# How does regularization prevent overfitting?




$$g(z) = \tanh(z)$$

$$\lambda \uparrow$$

$$W^{[L]} \downarrow$$

$$z^{[L]} = W^{[L]} a^{[L-1]} + b^{[L]}$$

Every layer  $\approx$  linear.



$$J(\dots) = \underbrace{\sum_i \mathcal{L}(\hat{y}^{(i)}, y^{(i)})}_{\text{training loss}} + \underbrace{\frac{\lambda}{2m} \sum_L \|W^{[L]}\|_F^2}_{\text{regularization term}}$$






deeplearning.ai

# Regularizing your neural network

---

## Dropout regularization

# Dropout regularization



↑  
0.5    ↑  
0.5    ↑  
0.5



# Implementing dropout (~~Inverted dropout~~)

Illustrate with layer  $l=3$ .  $keep-prob = 0.8$   $0.2$

$$\rightarrow \boxed{d3} = \text{np.random.rand}(a3.shape[0], a3.shape[1]) < \text{keep-prob}$$

$$\underline{a3} = \text{np.multiply}(a3, d3)$$

#  $a3 \neq d3$ .

$$\rightarrow \boxed{a3 /= \text{keep-prob}} \leftarrow$$

50 units.  $\leadsto$  10 units shut off

$$z^{[4]} = w^{[4]} \cdot \underline{a^{[3]}} + b^{[4]}$$

$\uparrow$

reduced by 20%.

$$/= \underline{0.8}$$

Test

# Making predictions at test time

$$a^{[0]} = X$$

No drop out.

$$z^{[1]} = W^{[1]} a^{[0]} + b^{[1]}$$

$$a^{[1]} = g^{[1]}(z^{[1]})$$

$$z^{[2]} = W^{[2]} \underline{a^{[1]}} + b^{[2]}$$

$$a^{[2]} = \dots$$

↓  
↑  
y

/= keep-prob



deeplearning.ai



# Regularizing your neural network

---

## Understanding dropout

# Why does drop-out work?

Intuition: Can't rely on any one feature, so have to spread out weights.  $\leadsto$  Shrink weights.  $b_2$





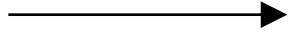
deeplearning.ai

# Regularizing your neural network



---

## Other regularization methods

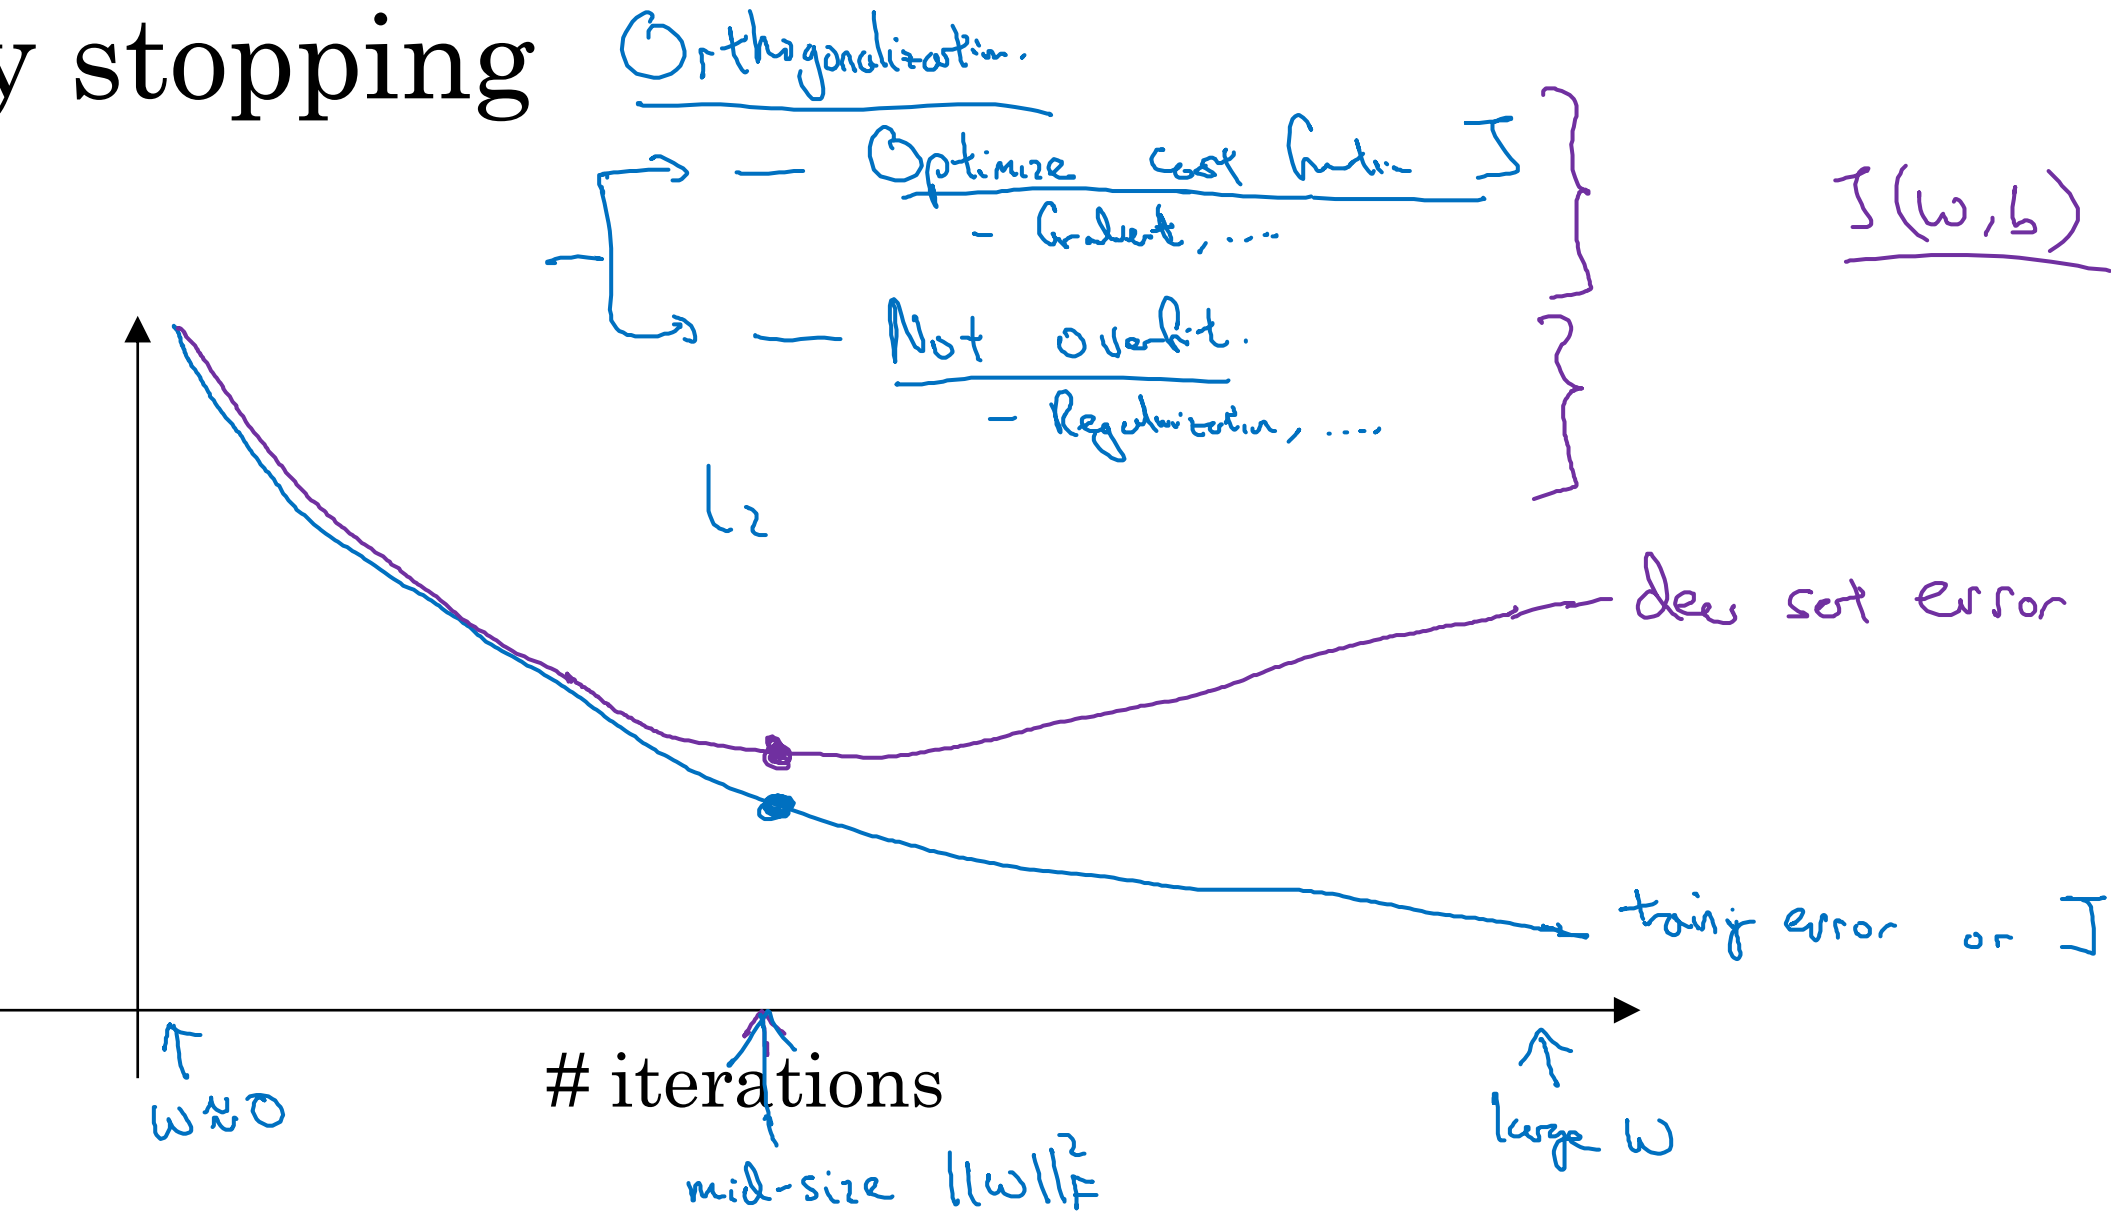
# Data augmentation



4



# Early stopping





deeplearning.ai

Setting up your  
optimization problem


---

Normalizing inputs



# Normalizing training sets

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$



Subtract mean:

$$\bar{\mu} = \frac{1}{n} \sum_{i=1}^n x^{(i)}$$

$$x := x - \mu$$

Normalize variance

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n x^{(i)} * x^{(i)T}$$

← element-wise



$$x /= \sigma^2$$

Use same  $\mu$   $\sigma^2$  to normalize test set.

# Why normalize inputs?



$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Unnormalized:



$x_1: 0 \dots 1$   
 $x_2: -1 \dots 1$   
 $x_3: 1 \dots 2$

Normalized:





deeplearning.ai


Setting up your  
optimization problem

---

Vanishing/exploding  
gradients

# Vanishing/exploding gradients

$L=150$



$g(z) = z$        $b^{[L]} = 0$

$\hat{y} = W^{[L]} \left( W^{[L-1]} W^{[L-2]} \dots W^{[2]} W^{[1]} x \right)$

$1.5^L$   
 $0.5^L$

$W^{[1]} > I$

$W^{[2]} < I$        $\begin{bmatrix} 0.9 & \\ & 0.9 \end{bmatrix}$


$W^{[2]} = \begin{bmatrix} 1.5 & 0 \\ 0 & 0.5 \end{bmatrix}$

$z^{[1]} = W^{[1]} x$   
 $a^{[1]} = g(z^{[1]}) = z^{[1]}$   
 $a^{[2]} = g(z^{[2]}) = g(W^{[2]} a^{[1]})$

$\hat{y} = W^{[L]} \left[ \begin{bmatrix} 1.5 & 0 \\ 0 & 0.5 \end{bmatrix}^{L-1} x \right]$

$1.5^{L-1} x$   
 $0.5^{L-1} x$

# Single neuron example



$$z = w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

large  $n \rightarrow$  Smaller  $w_i$

$$\text{Var}(w_i) = \frac{1}{n} \frac{2}{n}$$

$$\underline{w^{[1]}} = \text{np.random.randn}(\text{shape}) * \text{np.sqrt}\left(\frac{2}{n^{[1-1]}}\right)$$

ReLU  $g^{[2]}(z) = \text{ReLU}(z)$

Other variants:

tanh

$$\frac{1}{n^{[l-1]}}$$

Xavier initialization ↑

$$\sqrt{\frac{2}{n^{[l-1]} + n^{[1]}}}$$

↑



deeplearning.ai


Setting up your  
optimization problem

---

Numerical approximation  
of gradients

# Checking your derivative computation

I  $f(\theta) = \theta^3$   
 $\theta \in \mathbb{R}.$



$$g(\theta) = \frac{d}{d\theta} f(\theta) = f'(\theta)$$

$g(\theta) = 3\theta^2$

$\frac{dw}{db}$

$g(\theta) = 3 \cdot (1)^2 = 3$   
 when  $\theta = 1$

$$\frac{f(\theta + \epsilon) - f(\theta)}{\epsilon} \approx g(\theta)$$


$$\frac{(1.01)^3 - 1^3}{0.01} = 3.0301 \approx 3$$

0.0301  
 3.1  
 3.2

$\theta = 1$   
 $\theta + \epsilon = 1.01$   
 $\epsilon = 0.01$

# Checking your derivative computation

$$\underline{f(\theta) = \theta^3}$$



$$\left[ \frac{f(\theta + \epsilon) - f(\theta - \epsilon)}{2\epsilon} \approx \underline{g(\theta)} \right]$$

$$\frac{(1.01)^3 - (0.99)^3}{2(0.01)} = 3.0001 \approx 3$$

$$g(\theta) = 3\theta^2 = 3$$

approx error: 0.0001

(prev slide: 3.0301. error: 0.03)

$$\left\{ \begin{array}{l} f'(\theta) = \lim_{\epsilon \rightarrow 0} \frac{f(\theta + \epsilon) - f(\theta - \epsilon)}{2\epsilon} \quad \begin{array}{l} O(\epsilon^2) \\ 0.01 \\ \underline{0.0001} \end{array} \quad \left| \quad \frac{f(\theta + \epsilon) - f(\theta)}{\epsilon} \quad \begin{array}{l} \text{error: } O(\epsilon) \\ 0.01 \end{array} \end{array} \right.$$





deeplearning.ai

Setting up your  
optimization problem

---

**Gradient Checking**

# Gradient check for a neural network

Take  $W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]}$  and reshape into a big vector  $\theta$ .

$$J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = J(\theta)$$

Take  $dW^{[1]}, db^{[1]}, \dots, dW^{[L]}, db^{[L]}$  and reshape into a big vector  $d\theta$ .

Is  $d\theta$  the gradient of  $J(\theta)$ ?

# Gradient checking (Grad check)

$$J(\theta) = J(\theta_1, \theta_2, \theta_3, \dots)$$

for each  $i$ :

$$\rightarrow \underline{d\theta_{\text{approx}}[i]} = \frac{J(\theta_1, \theta_2, \dots, \overset{\downarrow}{\theta_i + \epsilon}, \dots) - J(\theta_1, \theta_2, \dots, \overset{\downarrow}{\theta_i - \epsilon}, \dots)}{2\epsilon}$$

$$\approx \underline{d\theta[i]} = \frac{\partial J}{\partial \theta_i} \quad | \quad d\theta_{\text{approx}} \approx d\theta$$

Checks

$$\rightarrow \frac{\|d\theta_{\text{approx}} - d\theta\|_2}{\|d\theta_{\text{approx}}\|_2 + \|d\theta\|_2}$$
$$\underline{\epsilon = 10^{-7}}$$

$$\approx \frac{10^{-7}}{10^{-5}} - \text{great!} \leftarrow$$
$$\rightarrow 10^{-3} - \text{worry.} \leftarrow$$



deeplearning.ai

Setting up your  
optimization problem

---

Gradient Checking  
implementation notes

# Gradient checking implementation notes

- Don't use in training – only to debug

$$\frac{d\theta_{\text{approx}}[\tilde{i}]}{\uparrow \uparrow} \longleftrightarrow \frac{d\theta[\tilde{i}]}{\uparrow}$$

- If algorithm fails grad check, look at components to try to identify bug.

$$\frac{db^{[L]}}{\uparrow} \quad \frac{dW^{[L]}}{\uparrow}$$

- Remember regularization.

$$\underline{J(\theta)} = \frac{1}{n} \sum_i \ell(y^{(i)}, \hat{y}^{(i)}) + \frac{\lambda}{2n} \sum_l \|W^{[l]}\|_F^2$$

$d\theta = \text{gradient of } J \text{ wrt. } \theta$

- Doesn't work with dropout.

$$\underline{J} \quad \underline{\text{keep-prob} = 1.0}$$

- Run at random initialization; perhaps again after some training.

$$\underline{W, b \approx 0}$$